



The TLM 2.0 Mixed Endianness Example

James Aldis

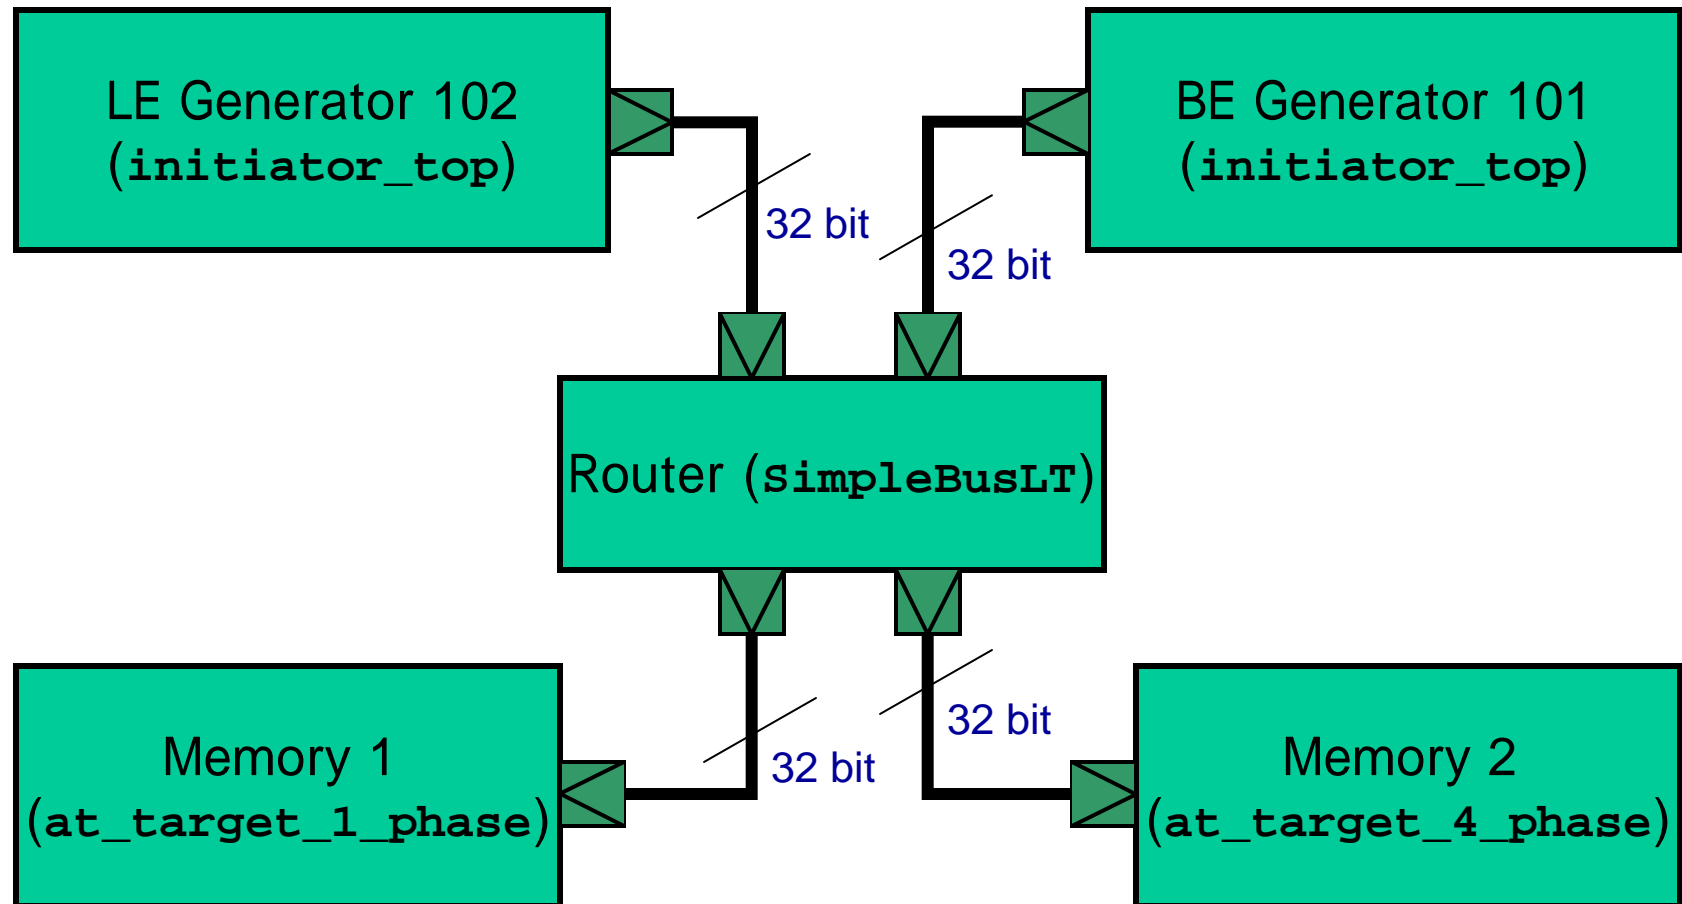
Texas Instruments France

January 2008

Modification of “It” Example

- Demonstrates the use of TLM endianness-conversion functions
- Illustrates data and address modification effects of BIG- and LITTLE-endian initiators sharing a memory
- Allows interactive experimentation
- Interactively execute “instructions” alternately on two initiators
 - one big- and one little-endian
 - store data from BE and view it from LE or vice-versa
 - store 32-bit data and view as 8-bit or 16-bit, etc

Platform Structure



tlm_initiator_socket



tlm_target_socket

How to run this example (Linux)

- Set `SYSTEMC_HOME`
- `cd examples/tlm/lt_mixed_endian/build-unix`
- `make`
- `make run` (uses default input)
- `./lt.exe` (interactive input)

How to run this example (MSVC)

- Open a explorer window on `examples/tlm/lt_mixed_endian/build-windows`
- Launch `lt.sln`
- Select '**Property Manager**' from the '**View**' menu
- Under '**lt_extension_mandatory > Debug | Win32**' select '**systemc**'
- Select '**Properties**' from the '**View**' menu
- Select '**User Macros**' under '**Common Properties**'
- Update the '**SYSTEMC**' entry and apply
- Select '**Debugging**' under '**Configuration Properties**'

Click in the '**Command Arguments**' box and type in

The logo for SYSTEMC, featuring the word "SYSTEMC" in a sans-serif font, with a stylized blue and yellow swoosh graphic above the letters "Y" and "E".


SYSTEMC™

Component details

- **Generator 101**
 - 32-bit Big-endian initiator
 - stdin provides instructions to execute (loads and stores)
 - stdin provides data for stores
 - data from loads written to stdout
- **Generator 102**
 - Identical but Little-endian
- **Instruction Set for both Generators**
 - l8, l16, l32 (load byte, halfword or word)
 - s8, s16, s32 (store byte, halfword or word)
 - w (switch control to other generator)
 - q (quit)
 - see `examples/tlm/lt_mixed_endian/results/input.txt` for example instructions
- **Bus and Memory**
 - Exactly as in lt example
 - Some kind of 32-bit routing and memory system


The System Being Modelled

- Both memories are simple arrays of 4-byte words
- Neither the memories nor the interconnect is aware of the endianness of the transactions
- Neither the memories nor the interconnect does any data modification:
 - what the initiators put on the bus goes unchanged to the memory
- Memories are endianness-neutral
 - provide a consistent memory image to BE and to LE initiators
 - “What I write, I read back the same”
 - there are no attributes on the hardware bus except those shown opposite




The System Being Modelled

- “What the other writes, I may read back at a different address”
- Example 1: single byte access
 - BE initiator writes byte A5 to address $66 = 64 + 2$
 - ◆ see opposite for hardware bus signals
 - LE initiator will find it at address $65 = 64 + 1$
 - ◆ see opposite for hardware bus signals



Address bus takes the index of the correct 32-bit word ($66/4 = 16$).

Fractional part of the address (2) used inside initiator to select a byte lane according to **big-endian** convention (more significant bits are at lower addresses)



Little-endian convention is that more significant bits are at *higher* addresses. To get the same data, fractional part of the address is 1

The System Being Modelled

- “What the other writes, I may read back at a different address”

- Same effect for aligned 16-bit or 32-bit data

- But no address change for 32-bit

- Example 2: 32-bit integer access

- BE initiator writes 32-bit integer 04030201 to address 256

- ♦ see opposite for hardware bus signals

- LE initiator will find the same data at the same address

- ♦ see opposite for hardware bus signals

address

0 00 40

bit index: 8 4 0

data in

04 03 02 01

24 16 8 0

byte enable

1 1 1 1

3 0

Address bus takes the index of the correct 32-bit word ($256/4 = 64$).

Big-endian convention is that more significant bits of the integer are at lower addresses, which are at more significant bits of the bus data word.

address

0 00 40

bit index: 8 4 0

data out

04 03 02 01

24 16 8 0

byte enable

1 1 1 1


3 0

Little-endian convention is that more significant bits of the integer are at *higher* addresses, which are at more significant bits of the bus data word. Therefore it is the same as big-endian!


The System Being Modelled

- “What the other writes, I may read back distorted”
- Example 3: write 4 consecutive bytes and read back as an integer

- BE initiator writes bytes 04, 03, 02, 01 to addresses 0, 1, 2, 3
 - ♦ see opposite for hardware bus signals
- Both initiators will read the integer 04030201 at address 0
- LE initiator writes bytes 04, 03, 02, 01 to addresses 0, 1, 2, 3
 - ♦ see opposite for hardware bus signals
- Both initiators will read the integer 01020304 at address 0



Big-endian convention is that more significant bits of the data bus word are lower addresses.



Little-endian convention is that more significant bits of the data bus word are *higher* addresses.

The TLM Model of the System

- The TLM model correctly models the above data and address distortions and all other possible ones
 - In particular it gets hairy for non-address-aligned transactions
- The address, data array and byte enable array in the transaction payload object are
 - identical to the internal opcode of the initiator
 - ♦ *if the initiator endianness matches the host CPU endianness*
 - modified
 - ♦ *if the initiator endianness is different from the host CPU's*
- The internal data storage of the memory models
 - is not visible (we are not using DMI in this example)
 - But could be a simple memcpy() between the data array in the transaction payload object and an array of unsigned char in the memory model
- Therefore we can say that
 - *the TLM 2.0 interfaces are always “host-endian”*
- Model is functionally identical on BE and LE host CPUs
 - *but internal structure (length, address, byte enables) of transaction payload objects will differ*